Figure 7 -- A0 vs A0': coverage went down, accuracy went up, because what was seen could be seen again

ANALYTIC ENTAILMENT -- this contrast DEMONSTRATES the mechanism, it does not TEST it. "n = 1 per arm, on worlds built by the same instance that theorized them" covers only sampling error. The objection that bites is analytic: A0''s toggle was designed so that every direction-by-polarity combination would have its own witness, and the adjudication rule -- admit a generalisation iff every case is witnessed -- then mechanically admits what it mechanically rejected in A0. The outcome follows from the construction; nothing was learned that was not built in. (papers/phase1-workshop/sections/03_a0.md 3.3.)

TWO VARIABLES, NOT ONE, AND NOT THE ONLY TWO -- the advertised change is latch -> toggle, and the explorer is THEN weakened on purpose. "Identical except" would be a false description and is not used: see panel C for all eight differences.

A. coverage went DOWN, accuracy went UP

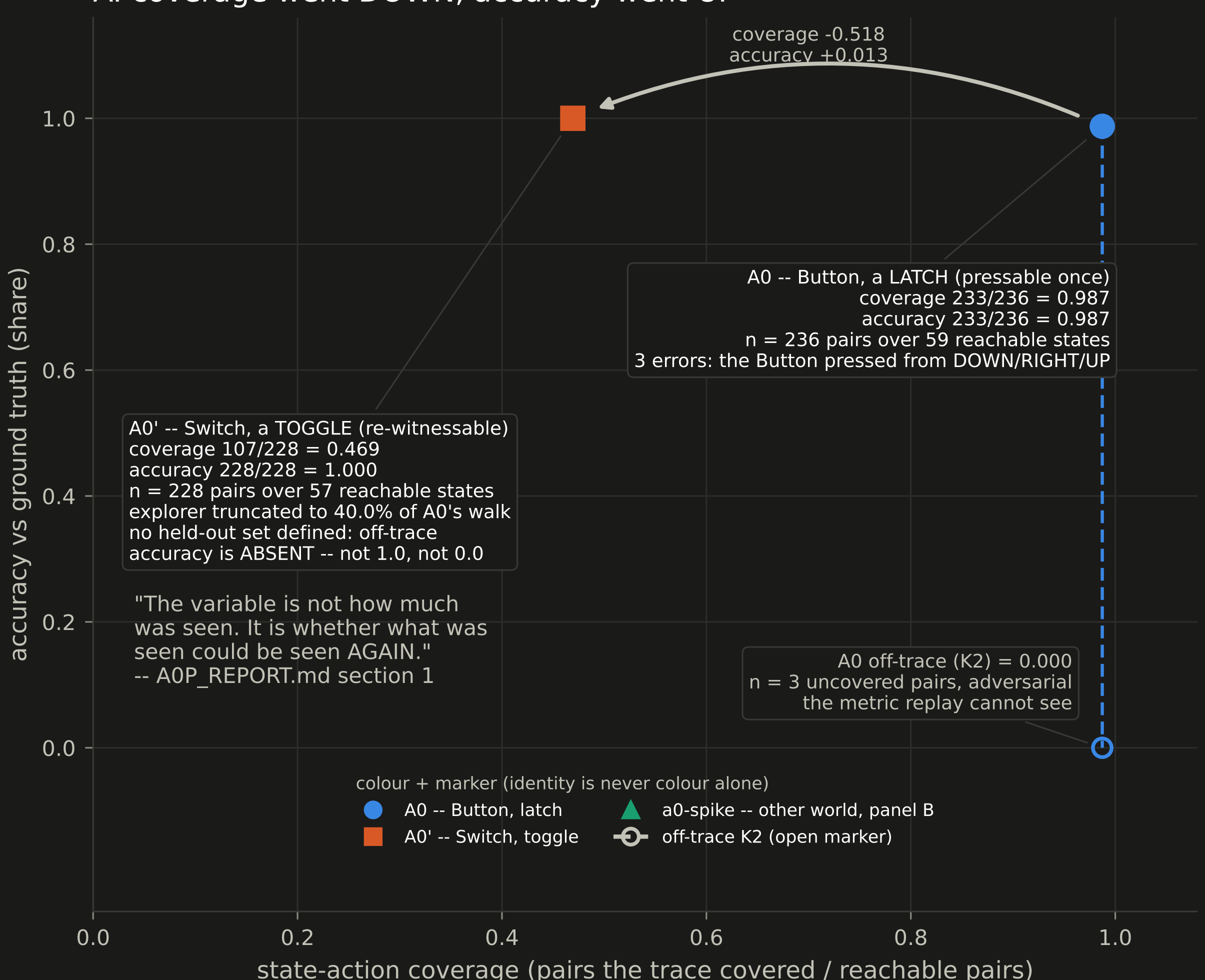

B. 0.000 beside 1.000 is meaningless without the denominators

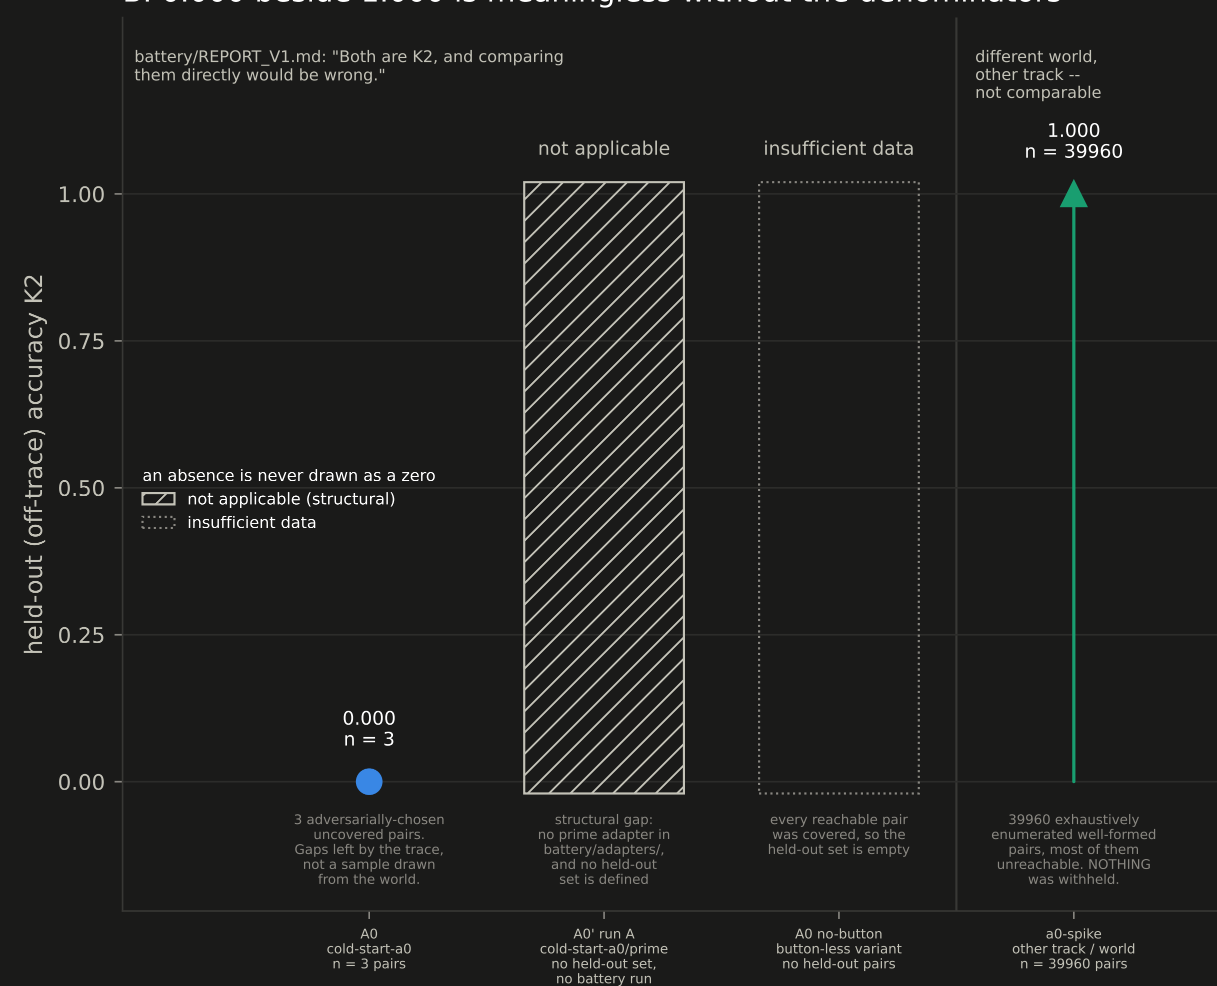

C. "identical except" is a false description and is not used: shaded = the two ADVERTISED variables, the six below differ too

quantity	A0 (cold-start-a0)	A0' (cold-start-a0/prime)
mechanism [ADVERTISED VARIABLE]	Button, a LATCH -- press witnessed in 1 of 4 directions	Switch, a TOGGLE -- witnessed in 8 of 8 direction x polarity
explorer [ADVERTISED VARIABLE]	exhaustive: 275 steps, 233/236 pairs covered	truncated: 110-transition budget = 40.0% of A0's walk, 107/228 covered
reachable states	59	57
state-action pairs	236	228
rules in the manual	7	21
objects in the manual	3 (prose only: 03_a0.md 3.3; no declared source)	3 (engines.tracks, after re-identification)
executable probes emitted	NOT READABLE from a declared source (A0P_REPORT section 1: 0 of 22 designed)	13 of 27 designed
revisions	0 (nothing came back from certify)	0 in run A, 1 in run B (seeded, panel D)

D. A0' run B -- a controlled experiment with a seeded error, not a discovery: one false clause, invisible to replay

full-history replay (cheap certify)	GREEN -- AND BLIND	the seeded clause never fires in the 110-transition history, so replay cannot see it. Exactly as predicted.
Lean transcription	CAUGHT IT	ArenaEscape: step sends the mover to (2, 4), which the board does not list as arena (from (2, 3) on right)
coverage probe (1 probe run)	CAUGHT IT	push_onto_crate has 2 firing states and neither is in the trace -- navigate, write the prediction down first, execute, world disagrees: refuted.

repair: delete the clause -- 1 revision, and it was a deletion. accuracy 0.991228 (226/228) -> 1.000000 (228/228). No multi-round repair is exercised anywhere in either run.

A0' is cold-start-a0/prime, not a0-spike: PLAN.md section 2 named a0-spike as A0', and a0-spike is a separate cold start on a different world run by the other track (03_a0.md 3.5), so it appears only in panel B, behind a divider. | The two K2 values in panel B are not comparable and are never shown without their denominators: A0's n = 3 is the adversarial gaps the trace left, a0-spike's n = 39960 is an exhaustive enumeration with nothing withheld. battery/REPORT_V1.md: "Both are K2, and comparing them directly would be wrong." battery/REPORT_V2.md: "Comparability was bought. Safety was not." | A0' has no held-out set, so its K2 is absent, not zero; and A0' has no battery run at all -- battery/adapters/ carries a0.py, a0_spike.py, a2.py and no prime adapter -- so the battery panel cannot include it, and the gap is drawn as a gap. | Both worlds are self-built: one instance built the world and adjudicated it (A0P_REPORT.md 5.5, "the seal has the same hole as A0's"), and the manuals differ in eight ways, not one (panel C). | A0's executable-probe count and both object counts are not readable from any source declared in figures/sources.py; they are labelled on the plate rather than drawn as numbers this pipeline can hash.