



Figure 4 -- A3 migration: carrying the book vs starting again (C3: a domain written once is worth carrying)


A. same level, books or no books: the like-for-like bill (n = 1 run per arm, no variance, no error bars)




B. cost to first plan (same two arms; I1 is a different level)




C. the one row with a real sample size



D. the safety valve: static caught neither, replay caught both



E. what the transfer arm was handed vs read off the frame



panels A, B, C -- the arms

- I2_from_scratch -- the control: level 2 from nothing, carrying nothing
- I2_transfer -- carries domain.dsl + playbook.dsl from level 1, unchanged
- * toolchain conformance, not world-facing work (panel A)
- I1_cold_start -- a DIFFERENT LEVEL; drawn only on panel C, never in a ratio

panels D, E -- absence encodings; an absent value is never drawn as a zero

- ▨ not applicable (structural)
- ⋯ insufficient data

n = 1 RUN PER ARM. Nothing here is replicated and no artefact carries a variance, so no error bar is drawn and none could be; the only genuine sample size on the plate is panel C's n = 248/252/252 reachable (state, action) pairs, printed beside each 1.0000. * THE TOOLCHAIN TAX: 2 of the control arm's 5 theorize rounds went to toolchain conformance (grounding lifted rules, then object names) rather than to the world -- 40% of its adjudication budget -- so the theorize_rounds saving overstates transfer by 2 rounds. THE BOTTOM THREE LINES SAVE NOTHING (1:1, 3:3, 1:1) and are drawn for that reason: a table showing savings there would be measuring something other than transfer. WHO CARRIED THE BOOKS: the person who carried them to level 2 is the person who wrote them for level 1, and they already knew the answer. The control arm's blind was PARTIALLY BROKEN and recorded as incident A3-I1 -- object names and the law's name are contaminated -- so no cross-arm agreement on naming is claimed anywhere on this figure. LEVELS, NOT GAMES: the two levels share a mechanism set by construction; A3 says nothing about carrying a domain between games with different mechanics. THE BILL IS STRUCTURAL, NOT ECONOMIC: it counts frames, actions, engine stages, candidates, rounds and clauses, and does NOT count model calls -- the largest term in a real bill. The zeros are real and the right shape; converting them to dollars is not something this experiment did. I1_cold_start appears only in panel C, labelled by level, and in the CSV. Ratios against I1_cold_start compare across levels and therefore mix in whatever differs between them. The like-for-like block is the same level with and without the books, and is the number to quote. A 0 on panel A or B means the arm genuinely did none of that (a real zero); the hatched and dotted cells in panels D and E are values that are absent, and absence is never drawn as a zero.