



Figure 2 -- bill shape: the per-turn cost curve (C2: bought early, spent late)

A. absolute: cumulative model cost




B. normalised: where E2 and E3 are read



C. the confound first: environment actions per turn, failures in red (both published failure bands are in the caveat)



D. context growth: transcript or manual?



E4 IS THE BATTERY'S NUMBER, NOT THIS PLATE'S. Higher means context is accelerating: R2 of a quadratic fit to context tokens per turn, minus R2 of a linear one. It reads the token series, not the priced one, so it survives a change in the price list.

ABSENT, NOT ZERO -- 40 of 110 drawn runs carry no E4:
 8 run(s):
 fewer than five turns; a quadratic needs room
 32 run(s):
 no battery run at all; battery v2 scores the arms
 bare_cc, schema_repro, theoria_a0, theoria_a0_spike, theoria_a2

THE TWO ARMS ARE NOT THE BATTERY'S NUMBER, BUT THE PLATE'S. The battery's own floor is quoted rather than restated -- it reports those runs as insufficient-data with the reason "fewer than 8 turns; a short run is trivially front-loaded" -- so a short run gets that verdict rather than a flattening number. C2 SHAPE, MEASURED AND NEGATIVE. Theoria.md 1.6 predicts a bill that is front-heavy and then tapers to nothing as the theory converges. Over the 4 drawn theoria leg(s) long enough to have a shape (at least 4 distinct billed steps), 1 is/are front-heavy and 1 taper(s) to under 10% of peak: 20260728T083400Z-E3-sk48-carried-v2 7 steps, front-half share 1.00, tail/peak 0.00; 20260731T1430Z-A3-level2-carried-r3 8 steps, front-half share 0.50, tail/peak 0.88; 20260731T1500Z-A3-sk48-carried-l1 4 steps, front-half share 0.44, tail/peak 0.82; 20260801T001851Z-R1b-g50t-a 6 steps, front-half share 0.46, tail/peak 0.64. THE PLATE SHOWS NO CONVERGENCE. It is written here rather than left to the reader because a flat curve and a converging one look alike at a glance when the curve is six points long. The taper that does appear is on 20260728T083400Z-E3-sk48-carried-v2, and that leg's own manifest records its outcome as budget_exhausted -- its zero-cost steps are the run still acting after the money stopped. The desk stopped being called because the budget ended, not because the manual stopped being surprised. A budget cutoff drawn on a cost axis is indistinguishable from convergence, and it is not convergence. Read against how the legs ended (budget_exhausted: 1; spend_gate_triggered: 3), not one of them stopped because it ran out of surprises. C2 is therefore UNCONFIRMED on this evidence, and the reason is where the legs were cut off rather than what they showed before that.