



Figure 2 -- bill shape: the per-turn cost curve (C2: bought early, spent late)


A. absolute: cumulative model cost




B. normalised: where E2 and E3 are read



C. the confound first: environment actions per turn, failures in red (both published failure bands are in the caveat)



D. context growth: transcript or manual?



E4 IS THE BATTERY'S NUMBER, NOT THIS PLATE'S. Higher means context is accelerating: R2 of a quadratic fit to context tokens per turn, minus R2 of a linear one. It reads the token series, not the priced one, so it survives a change in the price list.

ABSENT, NOT ZERO -- 40 of 110 drawn runs carry no E4:
 8 run(s):
 fewer than five turns; a quadratic needs room
 32 run(s):
 no battery run at all; battery v2 scores the arms bare_cc, schema_repro, theoria_a0, theoria_a0_spike, theoria_a2

billed model calls E4 was fitted over (battery E4 support turns)

THE TWO ARMS ARE NOT THE SAME CHAIN -- bare_cc buys one model call per tick; a theoria turn buys a desk call that theorises across the whole run -- 9 calls covered 25 actions, so its curve is short and tall by construction and the vertical gap in panel A is NOT a like-for-like markup. Per successful action the comparison that does hold is USD 0.7100 (theoria, 25 actions) against USD 0.1459 (bare_cc opus, baseline-arms/BUDGET_REPORT.md 2.1, opus row) -- and even that is one theoria run against a pilot. Theoria dollars are the provider's own arithmetic; the repo price table disagrees by -1.5% and that is a finding about the table, not the run. The third arm is still absent, and absence is not zero: there is no Schema arm in this ledger (baseline-arms/SCHEMA_LOCATE.md), so the model ladder stands in for it (battery/DECISIONS.md D-B-004, weaker by REPORT_V0's own note). Cost is summed over retried calls sharing a turn. Panel C is the price-list confound: with a large share of steps failing, E5 cost-per-action tracks token pricing, not skill. BOTH FIGURES FOR THAT SHARE TRAVEL, because they disagree: battery/REPORT_V0.md says 27-45%, and papers/phase1-workshop/REVIEW.md recomputes 28.3-45.1% and records that the 27% lower bound does not reproduce. The panel is annotated with REPORT_V0's band, which is the one it was drawn against. The envelope campaign's 16 committed ledger shard(s) are folded in, declared by rule in figures/sources.py rather than by a list that ages. E2/E3/E4 ARE THE BATTERY'S NUMBERS, NOT THIS PLATE'S. They are read from battery/artifacts/capability_spectrum.json, whose definitions are: E2, Share of total cost spent in the first 25% of turns. High means front-loaded: the arm paid to understand, then coasted. E3, Fraction of the run's turns needed to reach 90% of its total cost. Low means the bill settled early. E4, R^2 of a quadratic fit to context tokens per turn minus R^2 of a linear fit. Positive means context is accelerating. Recomputing them here would be a second definition of a Phase 4 primary endpoint. THE TWO TURN AXES ARE CHECKED, NOT ASSUMED TO MATCH: the battery counts turns in model-call order (battery/INPUT_FORMAT.md gap 5) and this plate counts step_idx, so E3's crossing is marked only on runs where the two coincide. On this build they agree on all 67 run(s) that carry a battery turn count, and that agreement is the licence to draw the marks at all. 43 drawn run(s) carry no E2 at all, of which 32 have no battery run: the live theoria arm is in none of battery v2's arms (bare_cc, schema_repro, theoria_a0, theoria_a0_spike, theoria_a2), so its front-load index is ABSENT, not low. The battery's own floor is quoted rather than restated -- it reports those runs as insufficient-data with the reason "fewer than 8 turns; a short run is trivially front-loaded" -- so a short run gets that verdict rather than a flattering number. C2 SHAPE, MEASURED AND NEGATIVE. Theoria.md 1.6 predicts a bill that is front-heavy and then tapers to nothing as the theory converges. Over the 4 drawn theoria leg(s) long enough to have a shape (at least 4 distinct billed steps), 1 is/are front-heavy and 1 taper(s) to under 10% of peak: 20260728T083400Z-E3-sk48-carried-v2 7 steps, front-half share 1.00, tail/peak 0.00; 20260731T1430Z-A3-level2-carried-r3 8 steps, front-half share 0.50, tail/peak 0.88; 20260731T1500Z-A3-sk48-carried-l1 4 steps, front-half share 0.44, tail/peak 0.82; 20260801T001851Z-R1b-g50t-a 6 steps, front-half share 0.46, tail/peak 0.64. THE PLATE SHOWS NO CONVERGENCE. It is written here rather than left to the reader because a flat curve and a converging one look alike at a glance when the curve is six points long. The taper that does appear is on 20260728T083400Z-E3-sk48-carried-v2, and that leg's own manifest records its outcome as budget_exhausted -- its zero-cost steps are the run still acting after the money stopped. The desk stopped being called because the budget ended, not because the manual stopped being surprised. A budget cutoff drawn on a cost axis is indistinguishable from convergence, and it is not convergence. Read against how the legs ended (budget_exhausted: 1; spend_gate_tripped: 3), not one of them stopped because it ran out of surprises. C2 is therefore UNCONFIRMED on this evidence, and the reason is where the legs were cut off rather than what they showed before that.