



Figure 2 -- bill shape: the per-turn cost curve (C2: bought early, spent late)


A. absolute: cumulative model cost




B. normalised: where E2 and E3 are read



C. the confound first: environment actions per turn, failures in red (both published failure bands are in the caveat)



D. context growth: transcript or manual?



E4 IS THE BATTERY'S NUMBER, NOT THIS PLATE'S.
Higher means context is accelerating: R2 of a quadratic fit to context tokens per turn, minus R2 of a linear one. It reads the token series, not the priced one, so it survives a change in the price list.

ABSENT, NOT ZERO -- 33 of 103 drawn runs carry no E4:
8 run(s):
fewer than five turns; a quadratic needs room
25 run(s):
no battery run at all; battery v2 scores the arms
bare_cc, schema_repro, theoria_a0, theoria_a0_spike, theoria_a2

THE TWO ARMS ARE NOT PRICED IN THE SAME UNIT. A bare_cc turn buys one model call that picks one action; a theoria turn buys a desk call that theorises across the whole run -- 30 calls covered 30 actions, so its curve is short and tall by construction and the vertical gap in panel A is NOT a like-for-like markup. Per successful action the comparison that does hold is USD 0.2802 (theoria, 30 actions) against USD 0.1459 (bare_cc opus, baseline-arms/BUDGET_REPORT.md 2.1, opus row) -- and even that is one theoria run against a pilot. Theoria dollars are the provider's own arithmetic; the repo price table disagrees by -1.8% and that is a finding about the table, not the run. The third arm is still absent, and absence is not zero: there is no Schema arm in this ledger (baseline-arms/SCHEMA_LOCATE.md), so the model ladder stands in for it (battery/DECISIONS.md D-B-004, weaker by REPORT_V0's own note). Cost is summed over retried calls sharing a turn. Panel C is the price-list confound: with a large share of steps failing, E5 cost-per-action tracks token pricing, not skill. BOTH FIGURES FOR THAT SHARE TRAVEL, because they disagree: battery/REPORT_V0.md says 27-45%, and papers/phase1-workshop/REVIEW.md recomputes 28.3-45.1% and records that the 27% lower bound does not reproduce. The panel is annotated with REPORT_V0's band, which is the one it was drawn against. The envelope campaign's 16 committed ledger shard(s) are folded in, declared by rule in figures/sources.py rather than by a list that ages. E2/E3/E4 ARE THE BATTERY'S NUMBERS, NOT THIS PLATE'S. They are read from battery/artifacts/capability_spectrum.json, whose definitions are: E2, Share of total cost spent in the first 25% of turns. High means front-loaded: the arm paid to understand, then coasted. E3, Fraction of the run's turns needed to reach 90% of its total cost. Low means the bill settled early. E4, R^2 of a quadratic fit to context tokens per turn minus R^2 of a linear fit. Positive means context is accelerating. Recomputing them here would be a second definition of a Phase 4 primary endpoint. THE TWO TURN AXES ARE CHECKED, NOT ASSUMED TO MATCH: the battery counts turns in model-call order (battery/INPUT_FORMAT.md gap 5) and this plate counts step_idx, so E3's crossing is marked only on runs where the two coincide. On this build they agree on all 67 run(s) that carry a battery turn count, and that agreement is the licence to draw the marks at all. 36 drawn run(s) carry no E2 at all, of which 25 have no battery run: the live theoria arm is in none of battery v2's arms (bare_cc, schema_repro, theoria_a0, theoria_a0_spike, theoria_a2), so its front-load index is ABSENT, not low. The battery's own floor is quoted rather than restated -- it reports those runs as insufficient-data with the reason "fewer than 8 turns; a short run is trivially front-loaded" -- so a short run gets that verdict rather than a flattering number.